

SVM

Do you smell what's SVM is cooking?

Files: <http://bit.ly/2kOA6xY>

Alexandr Notchenko

Alexandr.Notchenko@skolkovotech.ru

Score board: <http://bit.ly/2jOAejv>

Plan

- Quickly re-iterating terminology and concepts about SVM from lecture.
- Checking out (L-)SVC and SVR functions in [sklearn.svm](https://scikit-learn.org/stable/modules/svm.html) , discussing their signatures and arguments
- Looking into what's cooking kaggle competition, Exploring dataset quickly
 - Thinking how can you convert text data into vectors for SVC input
 - discussing Using OneHotEncoding vs TfidfEncoder
 - 5-fold stratified validation split
 - Talking about properties of OneVsAll classification setup
 - Training final evaluating models
 - Plotting results
 - Changing hyper-parameters C (margin) and d (degree of polynomial) and plotting corresponding change in performance
- Finishing words

Reiterating terminology and concepts

SUPERVISED LEARNING

- **Training data:** sample S of size m drawn *i.i.d.* according to distribution D on $X \times Y$

$$S = \{(x_1, y_1), \dots, (x_m, y_m)\}$$

- **Problem** find hypothesis $h \in H$ with small generalization error
 - deterministic case: $y = f(x)$ is a deterministic function, only $x \sim D$
 - stochastic case: output is a probabilistic function of input, e.g. $y = f(x) + \varepsilon$

`sklearn.svm`: Support Vector Machines

The `sklearn.svm` module includes Support Vector Machine algorithms.

User guide: See the [Support Vector Machines](#) section for further details.

Estimators

<code>svm.SVC</code> ([C, kernel, degree, gamma, coef0, ...])	C-Support Vector Classification.
<code>svm.LinearSVC</code> ([penalty, loss, dual, tol, C, ...])	Linear Support Vector Classification.
<code>svm.NuSVC</code> ([nu, kernel, degree, gamma, ...])	Nu-Support Vector Classification.
<code>svm.SVR</code> ([kernel, degree, gamma, coef0, tol, ...])	Epsilon-Support Vector Regression.
<code>svm.LinearSVR</code> ([epsilon, tol, C, loss, ...])	Linear Support Vector Regression.
<code>svm.NuSVR</code> ([nu, C, kernel, degree, gamma, ...])	Nu Support Vector Regression.
<code>svm.OneClassSVM</code> ([kernel, degree, gamma, ...])	Unsupervised Outlier Detection.
<code>svm.l1_min_c</code> (X, y[, loss, fit_intercept, ...])	Return the lowest bound for C such that for C in (l1_min_C, infinity) the model is guaranteed not to be empty.

Low-level methods

<code>svm.libsvm.fit</code>	Train the model using libsvm (low-level method)
<code>svm.libsvm.decision_function</code>	Predict margin (libsvm name for this is <code>predict_values</code>)
<code>svm.libsvm.predict</code>	Predict target values of X given a model (low-level method)
<code>svm.libsvm.predict_proba</code>	Predict probabilities
<code>svm.libsvm.cross_validation</code>	Binding of the cross-validation routine (low-level routine)

```
class sklearn.svm. svc (C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False,
tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape=None,
random_state=None)
```

[source]

<http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC>

C : float, optional (default=1.0) - Penalty parameter C of the error term

kernel : string, optional (default='rbf') - kernel type ('linear', 'poly', 'rbf', 'sigmoid', 'precomputed' or a callable)

degree : int, optional (default=3) - Degree of the polynomial kernel function ('poly')

gamma : float, optional (default='auto', =1/n_features) - Precision for 'rbf', or multiplier for 'poly' and 'sigmoid'

coef0 : float, optional (default = 0) - Constant terms for 'poly' and 'sigmoid' kernels

class_weight : {dict, 'balanced'}, optional - Set the parameter C of class i to class_weight[i]*C for SVC. If not given, all classes are supposed to have weight one. The "balanced" mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as $n_samples / (n_classes * np.bincount(y))$

decision_function_shape : 'ovo', 'ovr' or None, default=None - Whether to return a one-vs-rest ('ovr') decision function of shape (n_samples, n_classes) as all other classifiers, or the original one-vs-one ('ovo') decision function of libsvm which has shape (n_samples, n_classes * (n_classes - 1) / 2)

tol : float, optional (default=1e-3) - Tolerance for stopping criterion (affects solution accuracy)

```
class sklearn.svm. LinearSVC (penalty='l2', loss='squared_hinge', dual=True, tol=0.0001, C=1.0, multi_class='ovr',  
fit_intercept=True, intercept_scaling=1, class_weight=None, verbose=0, random_state=None, [source]  
max_iter=1000)
```

<http://scikit-learn.org/stable/modules/generated/sklearn.svm.LinearSVC.html#sklearn.svm.LinearSVC>

C : float, optional (default=1.0) - Penalty parameter C of the error term.

loss : string, 'hinge' or 'squared_hinge' (default='squared_hinge') - Specifies the loss function. 'hinge' is the standard SVM loss (used e.g. by the SVC class) while 'squared_hinge' is the square of the hinge loss.


penalty : string, 'l1' or 'l2' (default='l2') - Specifies the norm used in the penalization. The 'l2' penalty is the standard used in SVC. The 'l1' leads to coef_ vectors that are sparse.




tol : float, optional (default=1e-4) - Tolerance for stopping criteria.


multi_class: string, (default='ovr') : - Determines the multi-class strategy if y contains more than two classes.

class_weight : {dict, 'balanced'}, optional - Set the parameter C of class i to class_weight[i]*C for SVC.

One-vs-rest classifier



Class 1: 
Class 2: 
Class 3: 



What's cooking dataset

<https://www.kaggle.com/c/whats-cooking>



Completed • Knowledge • 1,388 teams

What's Cooking?

Wed 9 Sep 2015 – Sun 20 Dec 2015 (13 months ago)

Data Files

File Name	Available Formats
sample_submission.csv	.zip (25.80 kb)
test.json	.zip (425.52 kb)
train.json	.zip (1.76 mb)

```
{  
  "id": 24717,  
  "cuisine": "indian",  
  "ingredients": [  
    "tumeric",  
    "vegetable stock",  
    "tomatoes",  
    "garam masala",  
    "naan",  
    "red lentils",  
    "red chili peppers",  
    "onions",  
    "spinach",  
    "sweet potatoes"  
  ],  
}
```


How can we convert text data into vectors for SVC?

- Sequences of words?
- Interaction of words on feature level or before?
- Can Stemming help?
- N-grams?

One-hot encoding

$V = \{\text{zebra}, \text{horse}, \text{school}, \text{summer}\}$

$v(\text{zebra}) = [1, 0, 0, 0]$

$v(\text{horse}) = [0, 1, 0, 0]$

$v(\text{school}) = [0, 0, 1, 0]$

$v(\text{summer}) = [0, 0, 0, 1]$

(+) Pros:

Simplicity

(-) Cons:

One-hot encoding can be memory inefficient

Notion of word similarity is undefined with one-hot encoding

n-grams

[illegible]

TFIDF

For a term i in document j :

$$w_{i,j} = tf_{i,j} \times \log \left(\frac{N}{df_i} \right)$$

$tf_{i,j}$ = number of occurrences of i in j

df_i = number of documents containing i

N = total number of documents

	Document 1	Document 2	Document 3	Document 4	Document 5	Document 6	Document 7	Document 8
Term(s) 1	10	0	1	0	0	0	0	2
Term(s) 2	0	2	0	0	0	18	0	2
Term(s) 3	0	0	0	0	0	0	0	2
Term(s) 4	6	0	0	4	6	0	0	0
Term(s) 5	0	0	0	0	0	0	0	2
Term(s) 6	0	0	1	0	0	1	0	0
Term(s) 7	0	1	8	0	0	0	0	0
Term(s) 8	0	0	0	0	0	3	0	0

Word Vector (Passage Vector) ←

Document Vector ↑

```
class sklearn.feature_extraction.text. TfidfTransformer (norm='l2', use_idf=True, smooth_idf=True,  
sublinear_tf=False)
```

[\[source\]](#)

Sources*:

1. [NYU ML Lectures - Support Vector Machines](#)
2. [NICTA MLSS - Classification notebook](#)
3. [Advanced BioStatistics - 8366 section 6.4](#)
4. [Support Vector Machine Tutorial - Wu, Shih-Hung \(Ph.D\) - Dept. of CSIE, CYUT](#)
5. [Michal Tadeusiak's solution](#)
6. [SaquibAhmad solution to What's cooking Kaggle competition](#)

*(Notebooks and slides from which to take some code/pictures/equations)